



Advanced Statistical Programming with R

Some more details on location-scale regression

Hannes Riebl

April 21, 2021

Questions?

Create your GitLab account

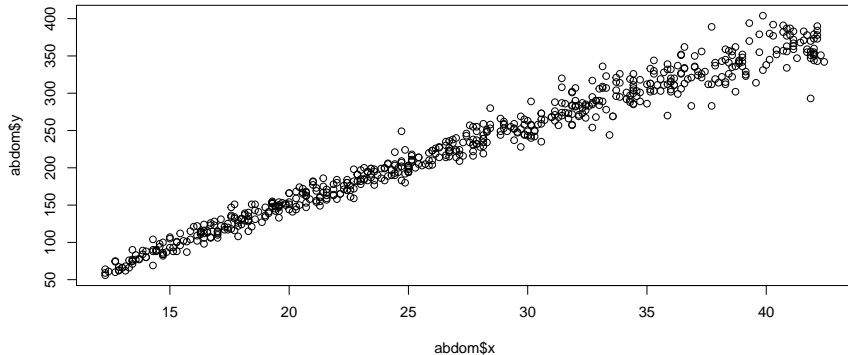


GitLab

Please log in to <https://gitlab.gwdg.de> **now!**

The abdom example

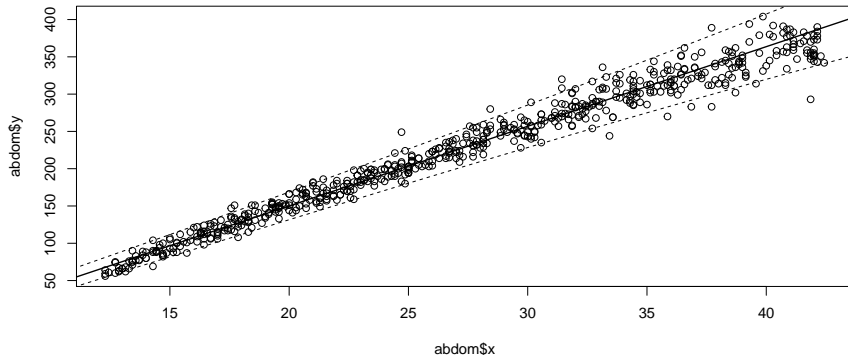
```
data(abdom, package = "gamlss.data")  
plot(abdom$x, abdom$y)
```



```
library(lslm)
mod <- lslm(location = y ~ x, scale = ~x, data = abdom)
summary(mod)
```

```
##
## Call:
## lslm(location = y ~ x, scale = ~x, data = abdom)
##
## Pearson residuals:
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max.
## -3.74000 -0.66540 -0.04960 -0.01809  0.64240  4.19900
##
## Location coefficients (identity link function):
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -63.47252    1.46549  -43.31  <2e-16 ***
## x            10.67805    0.06388  167.15  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Scale coefficients (log link function):
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.386819   0.096714   14.34  <2e-16 ***
## x            0.042992   0.003388   12.69  <2e-16 ***
```

```
plot(abdom$x, abdom$y)  
abline(coef(mod, "location")[1], coef(mod, "location")[2], lwd = 1.5)  
curve(f(x, "upper"), from = 5, to = 50, add = TRUE, lty = 2)  
curve(f(x, "lower"), from = 5, to = 50, add = TRUE, lty = 2)
```



Location-scale regression

The location-scale regression model is defined as

$$y_i \stackrel{\text{ind.}}{\sim} \mathcal{N}(x_i' \beta, (\exp(z_i' \gamma))^2),$$

where

- ▶ $i = 1, \dots, n$ is the observation index,
- ▶ y_i is the response variable,
- ▶ x_i is the vector of explanatory variables **for the mean**,
- ▶ β is the vector of regression coefficients **for the mean**,
- ▶ z_i' is the vector of explanatory variables **for the standard deviation**, and
- ▶ γ is the vector of regression coefficients **for the standard deviation**.

Maximum likelihood estimation

How can we estimate the full parameter vector of the location-scale regression model

$$\boldsymbol{\vartheta} = \begin{bmatrix} \boldsymbol{\beta} \\ \boldsymbol{\gamma} \end{bmatrix}$$

accurately, reliably, and efficiently?

Newton-Raphson / Fisher Scoring:

- In ML estimation, we are seeking the root of the score function $s(\vartheta)$.
- Basic idea of Newton-Raphson optimisation: Iteratively approximate $s(\vartheta)$ by a Taylor expansion (i.e. by a simple polynomial) and find the root of this approximation.
- First order Taylor approximation:

$$s(\vartheta) \approx s(\tilde{\vartheta}) - F(\tilde{\vartheta})(\vartheta - \tilde{\vartheta}).$$

- Setting the Taylor approximation to zero yields

$$\hat{\vartheta} = \tilde{\vartheta} + \frac{s(\tilde{\vartheta})}{F(\tilde{\vartheta})}.$$

- Newton-Raphson algorithm:
 - Determine a suitable starting value $\hat{\vartheta}^{(0)}$.
 - Update

$$\hat{\vartheta}^{(k+1)} = \hat{\vartheta}^{(k)} + \frac{s(\hat{\vartheta})^{(k)}}{F(\hat{\vartheta})^{(k)}}$$

until the change in the estimate becomes small.

- In statistics, the negative second derivative (i.e. the Fisher information) is often replaced by the expected Fisher information $F^*(\vartheta)$ (Fisher scoring).
 - $F^*(\vartheta)$ often has a simpler structure since terms with expectation zero cancel.
 - The solution should not be affected too much since

$$F(\vartheta) \approx F^*(\vartheta)$$

for large samples (and since the same score function is used).

- In certain classes of models, one can show that the expected Fisher information is ensured to be positive which is not always for the case the Fisher information $F(\vartheta)$.

Fisher scoring with a parameter **vector**

In the location-scale regression model, ϑ is a parameter **vector**, and hence the update step of the Fisher scoring algorithm is

$$\hat{\vartheta}^{(k+1)} = \hat{\vartheta}^{(k)} + \left(F^* \left(\hat{\vartheta}^{(k)} \right) \right)^{-1} s \left(\hat{\vartheta}^{(k)} \right).$$

Blockwise parameter updates

Until convergence, repeat:

- ▶ Update $\hat{\gamma}$.
- ▶ Update $\hat{\beta}$.

Why?

- ▶ The cross-covariance matrix of $s(\beta)$ and $s(\gamma)$ is $\mathbf{0}$.
- ▶ There is an analytic ML estimator for β given γ .

Estimating β given γ

Given γ , we know $\sigma_i = \exp(\mathbf{z}'_i \gamma)$ for all $i = 1, \dots, n$. Then we can estimate β with WLS:

$$\hat{\beta}^{\text{WLS}} = (\mathbf{X}'\mathbf{W}\mathbf{X})^{-1}\mathbf{X}'\mathbf{W}\mathbf{y},$$

where

$$\mathbf{W} = \begin{bmatrix} \sigma_1^{-2} & 0 & \dots & 0 \\ 0 & \sigma_2^{-2} & & \vdots \\ \vdots & & \ddots & 0 \\ 0 & \dots & 0 & \sigma_n^{-2} \end{bmatrix}.$$

The WLS estimator is equivalent to the ML estimator.

Fisher scoring for γ

We need the score $s(\gamma)$ and the expected Fisher information $F^*(\gamma)$ with respect to the parameter vector γ ...

Summing up the iterative ML estimation algorithm

Until convergence, repeat:

- ▶ Update $\hat{\gamma}$ with Fisher scoring, keeping $\hat{\beta}$ fixed at its current value.
- ▶ Update $\hat{\beta}$ with WLS, assuming the current value of $\hat{\gamma}$ is the true value.

Initializing $\hat{\beta}$ with OLS

We can initialize $\hat{\beta}$ with OLS:

$$\hat{\beta}^{\text{OLS}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}.$$

$\hat{\beta}^{\text{OLS}}$ is unbiased:

$$\mathbb{E}(\hat{\beta}^{\text{OLS}}) = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbb{E}(\mathbf{y}) = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{X}\beta = \beta.$$

Under mild regularity assumptions, $\hat{\beta}^{\text{OLS}}$ is also consistent.

However, $\hat{\beta}^{\text{OLS}}$ is no longer BLUE, because the homoskedasticity assumption of the Gauss-Markov theorem is violated.

Initializing $\hat{\gamma}$ with OLS

$$\sigma_i = \exp(\mathbf{z}_i' \gamma) \iff \log \sigma_i = \mathbf{z}_i' \gamma.$$

Hence, we are able to estimate γ with OLS if we can come up with a good estimator for $\log \sigma_i$.

The natural estimator for σ_i , $\hat{\sigma}_i = |y_i - \mu_i|$, is unbiased, but due to Jensen's inequality,

$$\mathbb{E}(\log \hat{\sigma}_i) < \log \mathbb{E}(\hat{\sigma}_i) = \log \sigma_i.$$

Deriving an unbiased estimator for $\log \sigma_i$

Is there a simple bias correction term?

Metropolis-Hastings-Algorithm:

- Define a starting value $\vartheta^{(0)}$.
- For $m = 1, \dots, M$ generate a proposal ϑ^* from a density

$$h(\vartheta^*|\vartheta^{(m-1)})$$

depending on the previous value $\vartheta^{(m-1)}$.

- Accept it with probability

$$\alpha(\vartheta^*|\vartheta^{(m-1)}) = \min \left\{ 1, \frac{f(\vartheta^*|\mathbf{x})}{f(\vartheta^{(m-1)}|\mathbf{x})} \frac{h(\vartheta^{(m-1)}|\vartheta^*)}{h(\vartheta^*|\vartheta^{(m-1)})} \right\}$$

If ϑ^* is accepted, set

$$\vartheta^{(m)} = \vartheta^*.$$

Otherwise set

$$\vartheta^{(m)} = \vartheta^{(m-1)}.$$

- Remarks:
 - The normalising constant in $f(\vartheta|\mathbf{x})$ cancels in the acceptance probability and is therefore not required.
 - The proposal density $h(\vartheta^*|\vartheta^{(m-1)})$ can be any density with the same support as the posterior $f(\vartheta|\mathbf{x})$.
 - If the proposal density is close to the posterior, the acceptance probability will always be close to 1.

- If the parameter vector $\boldsymbol{\vartheta}$ is of higher dimension, additional flexibility can be achieved by decomposing $\boldsymbol{\vartheta}$ as

$$\boldsymbol{\vartheta} = (\boldsymbol{\vartheta}'_1, \dots, \boldsymbol{\vartheta}'_j, \dots, \boldsymbol{\vartheta}'_k)'$$

Then, separate proposals for each subvector $\boldsymbol{\vartheta}_j$ are defined, i.e. a proposal $\boldsymbol{\vartheta}_j^*$ is generated from

$$h_j(\boldsymbol{\vartheta}_j^* | \boldsymbol{\vartheta}_1^{(m)}, \dots, \boldsymbol{\vartheta}_{j-1}^{(m)}, \boldsymbol{\vartheta}_j^{(m-1)}, \boldsymbol{\vartheta}_{j+1}^{(m-1)}, \dots, \boldsymbol{\vartheta}_k^{(m-1)})$$

and accepted with probability $\alpha(\boldsymbol{\vartheta}_j^*|\boldsymbol{\vartheta}_j^{(m-1)})$ defined as

$$\min \left\{ 1, \frac{f(\boldsymbol{\vartheta}_j^*|\boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x})}{f(\boldsymbol{\vartheta}_j^{(m-1)}|\boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x})} \frac{h_j(\boldsymbol{\vartheta}_j^{(m-1)}|\boldsymbol{\vartheta}_1^{(m)}, \dots, \boldsymbol{\vartheta}_{j-1}^{(m)}, \boldsymbol{\vartheta}_j^*, \boldsymbol{\vartheta}_{j+1}^{(m-1)}, \dots, \boldsymbol{\vartheta}_k^{(m-1)})}{h_j(\boldsymbol{\vartheta}_j^*|\boldsymbol{\vartheta}_1^{(m)}, \dots, \boldsymbol{\vartheta}_{j-1}^{(m)}, \boldsymbol{\vartheta}_j^{(m-1)}, \boldsymbol{\vartheta}_{j+1}^{(m-1)}, \dots, \boldsymbol{\vartheta}_k^{(m-1)})} \right\}$$

where

$$\boldsymbol{\vartheta}_{-j}^{(m)} = (\boldsymbol{\vartheta}_1^{(m)}, \dots, \boldsymbol{\vartheta}_{j-1}^{(m)}, \boldsymbol{\vartheta}_{j+1}^{(m-1)}, \dots, \boldsymbol{\vartheta}_k^{(m-1)})'$$

and

$$f(\boldsymbol{\vartheta}_j|\boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x})$$

is the full conditional of $\boldsymbol{\vartheta}_j$ (given all other parameters and the data \boldsymbol{x}).

- The full conditional $f(\boldsymbol{\vartheta}_j|\boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x})$ is proportional to the complete posterior, i.e. it can be determined by considering only those factors in $f(\boldsymbol{\vartheta}|\boldsymbol{x})$ that depend on $\boldsymbol{\vartheta}_j$.

Choice of the Proposal Density:

- If the full conditional is a known distribution, we can directly generate a proposal from this distribution that will be accepted with probability 1 (this is called a Gibbs update).
- Random walk proposal:

$$\vartheta_j^* = \vartheta_j^{(m-1)} + u_j, \quad u_j \sim N(0, \tau^2)$$

where τ^2 is a tuning parameter that determines the acceptance probability.

- Locally quadratic approximation of the log-full conditional: Determine the mode

$$\mathbf{m}_j = \arg \max_{\boldsymbol{\vartheta}_j} \log \left(f(\boldsymbol{\vartheta}_j | \boldsymbol{\vartheta}_{-j}^{(m)}, \mathbf{x}) \right)$$

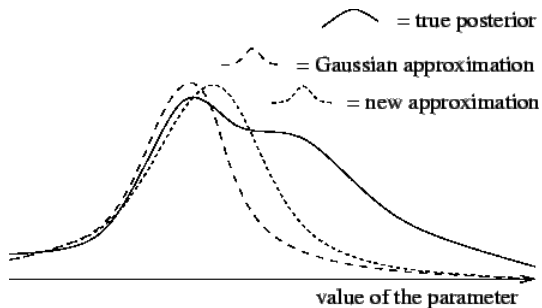
and the curvature at the mode

$$\mathbf{F}_j = - \frac{\partial^2}{\partial \boldsymbol{\vartheta}_j \partial \boldsymbol{\vartheta}_j'} \log \left(f(\boldsymbol{\vartheta}_j | \boldsymbol{\vartheta}_{-j}^{(m)}, \mathbf{x}) \right) \Big|_{\boldsymbol{\vartheta}_j = \mathbf{m}_j}$$

and propose the new state from

$$\mathrm{N}(\mathbf{m}_j, \mathbf{F}_j^{-1}).$$

Approximation of the posterior



Taken from: http://users.ics.aalto.fi/harri/thesis/valpola_thesis/node50.html

Proposal density in the `lslm` package

The `lslm` package uses the proposal density

$$\mathcal{N}\left(\boldsymbol{m}_j, \left(F^*\left(\boldsymbol{\vartheta}_j^{(m-1)} \mid \boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x}\right)\right)^{-1}\right),$$

where $\boldsymbol{m}_j = \boldsymbol{\vartheta}_j^{(m-1)} + \left(F^*\left(\boldsymbol{\vartheta}_j^{(m-1)} \mid \boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x}\right)\right)^{-1} \boldsymbol{s}\left(\boldsymbol{\vartheta}_j^{(m-1)} \mid \boldsymbol{\vartheta}_{-j}^{(m)}, \boldsymbol{x}\right)$.

Note that the `lslm` package assumes flat priors for β and γ , and hence, \boldsymbol{s} and F^* are simply the score and the expected Fisher information of the log-likelihood as before.

Next steps



1. Read the code of the `ls1m` package.
2. Read the book “R Packages” by Hadley Wickham and Jenny Bryan.
3. Read the book “Advanced R” by Hadley Wickham.
4. Get in touch with your topic supervisors.

Questions?